

SystemC/TLM: *AcceLera SystemC 2.3.1*

Franco Fummi



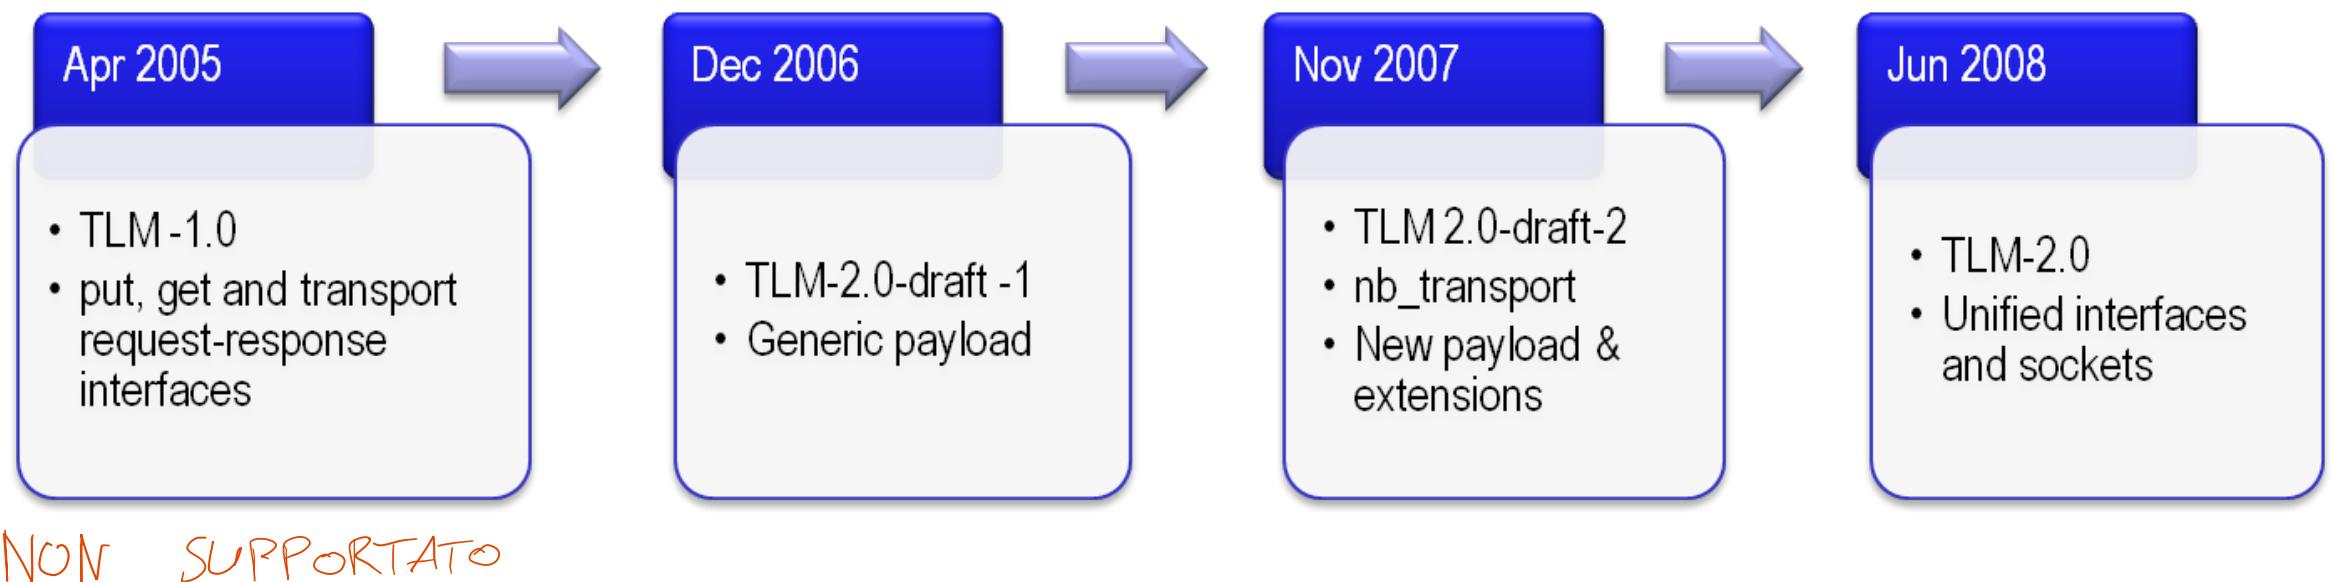
UNIVERSITÀ
di VERONA
Dipartimento
di **INFORMATICA**

Version 1.1

Contents

- Introduction
 - History
- Use cases, coding styles and mechanisms
 - Coding styles: Loosely vs. Approximately Timed
 - Initiators and targets
 - Blocking vs. Non-blocking transport
 - The time quantum
- Sockets and generic payload
- Verification principles

OSCI TLM Development



History of Releases (I)

- TLM-2005-04-08: TLM 1.0
 - Bidirectional interfaces
 - Transport, blocking
 - Unidirectional interfaces
 - FIFO, blocking, non-blocking
 - Separation between request and response
- TLM-2006-11-29: TLM 2.0 draft 1
 - TLM 1.0 + timing annotation
 - Generic payload + extensions

L'astrazione levels
TLM 1 → identificare una metodologia comune per tutte le committite
TLM 2 → modellare il sistema se possibile

History of Releases (II)

- TLM-2007-11-29: TLM 2.0 draft 2
 - Legacy support for TLM 1.0
 - Coding styles
 - Loosely-timed, Approximately-timed
 - New Transport interfaces
 - Temporal decoupling
 - Quantum keeper
 - Improved Debug and Direct Memory interfaces
- TLM-2008-06-09: TLM 2.0 standard
- TLM-2009: TLM 2.0.1 (minor changes and bug fixes)
- Standard migrated into ***Accellera SystemC 2.3.1***

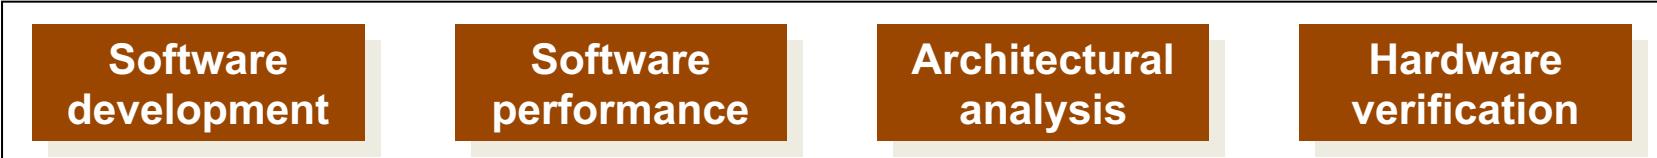
Use Cases, Coding Styles and Mechanisms

TLM : Transaction-Level Modeling

I tempi di simulazione sono molto più bassi, PERO' non so i cicli di clock esatti per la simulazione => l'unità di tempo diventa la transazione

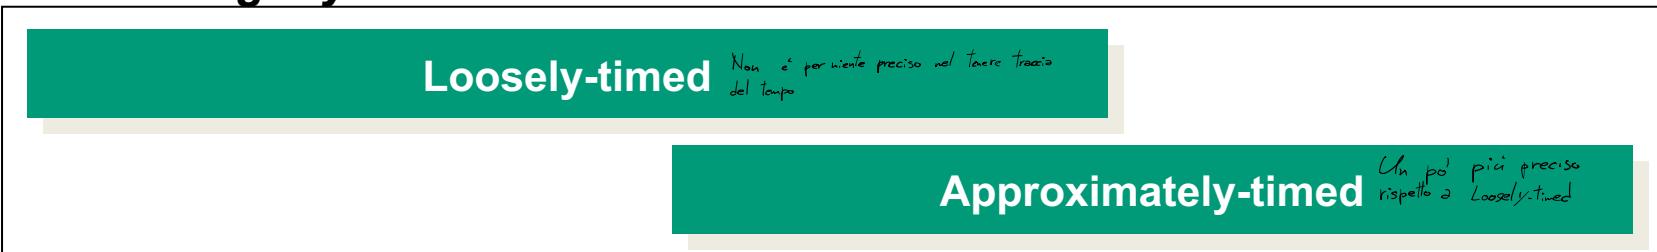
Chiedo l'esecuzione di un metodo, mando dati e ricevo altri di output

Use cases



I moduli possono essere esportati/
importati;

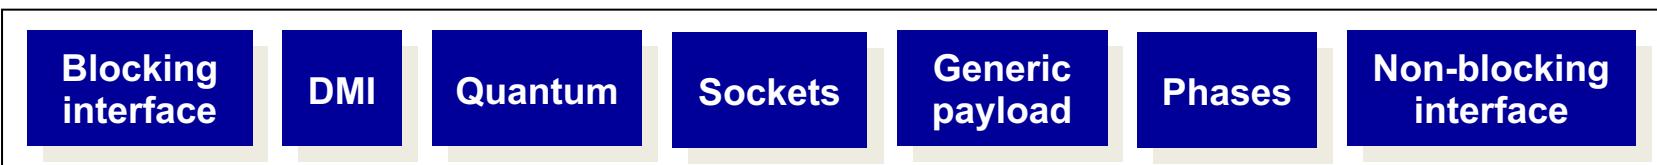
TLM-2 Coding styles



Se non ho bisogno di tempi esatti: dell'ho uso TLM che è molto veloce a
Simulare

Mechanisms

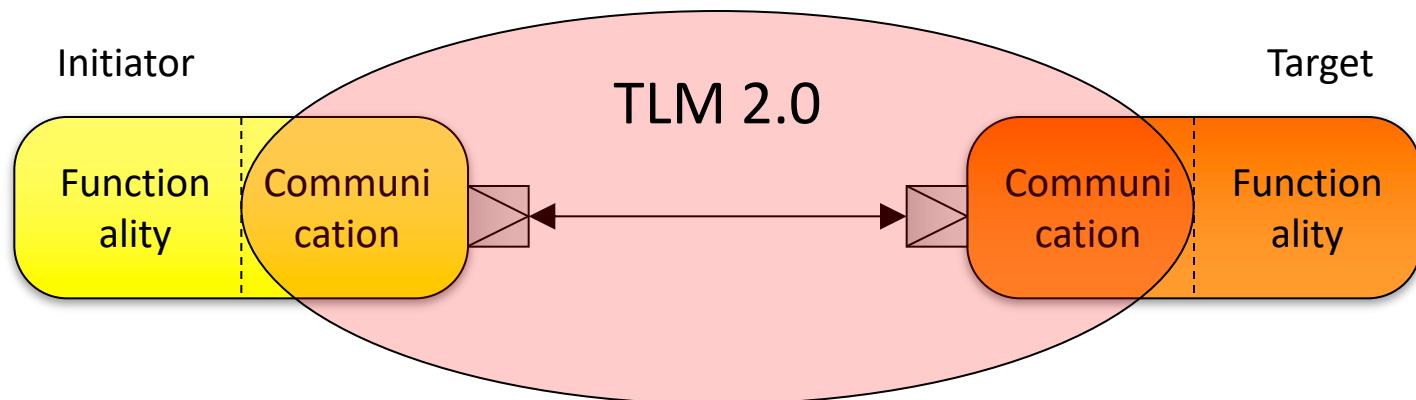
→ cosa la libreria TLM fornisce



Functionality vs Communication

- TLM manages to keep distinct functionality and communication in each module
- TLM 2.0 allows to model the communication part (i.e. how each module interacts with the others)

Le funzionalità possono essere in qualsiasi linguaggio, basta poterlo collegare a C++



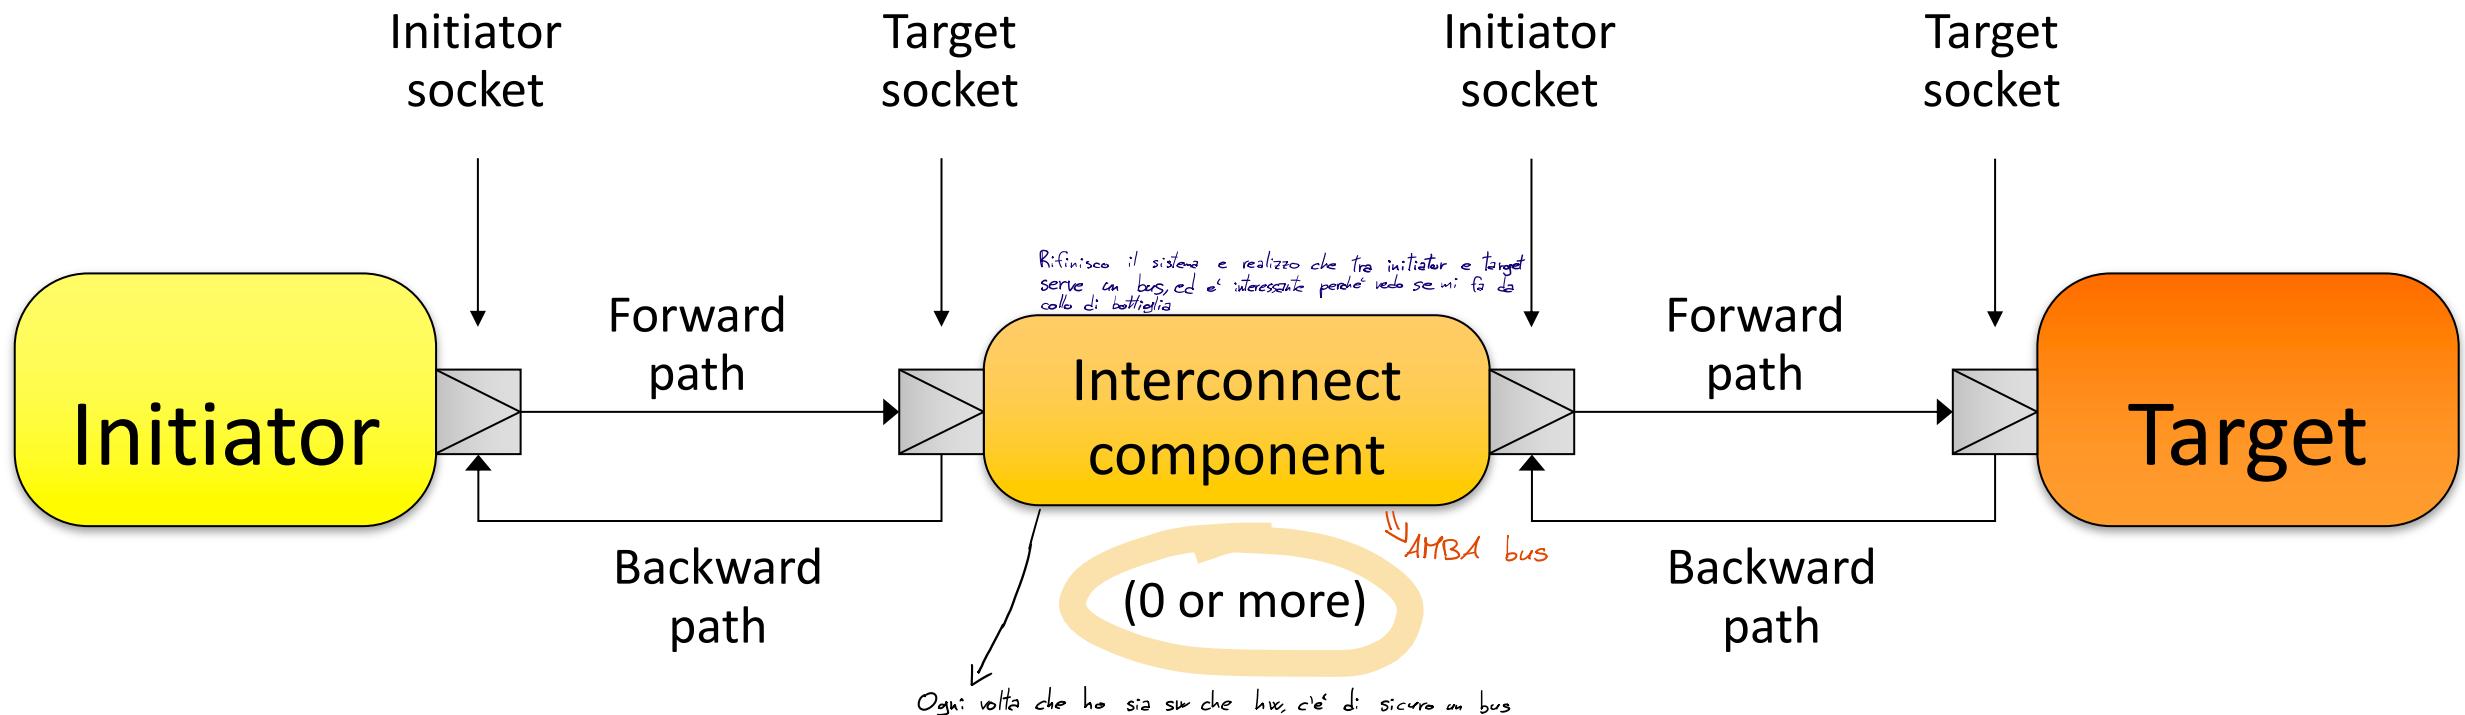
Transactions

- TLM relies on the notion of transaction
- A transaction consists of a data transfer from a design module to another one
 - Write and read operations are examples of transactions
 - It is usually represented by a generic payload object in the code
 - This object contains both data and control information
 - (e.g., address and type of command)
 - It is exchanged between modules through primitive calls

Socket, Initiator and Target (I)

- Communication is achieved by exchanging packets between an initiator module and a target module, through a socket
 - The initiator starts a transaction
 - The target is the end point of a transaction
 - An interconnect component is an intermediate point in the path from the initiator to the target
 - A socket connects two modules, and allows them to communicate by means of the available interfaces
 - The forward path runs from the initiator to the target
 - The backward path runs from the target to the initiator

Socket, Initiator and Target (II)



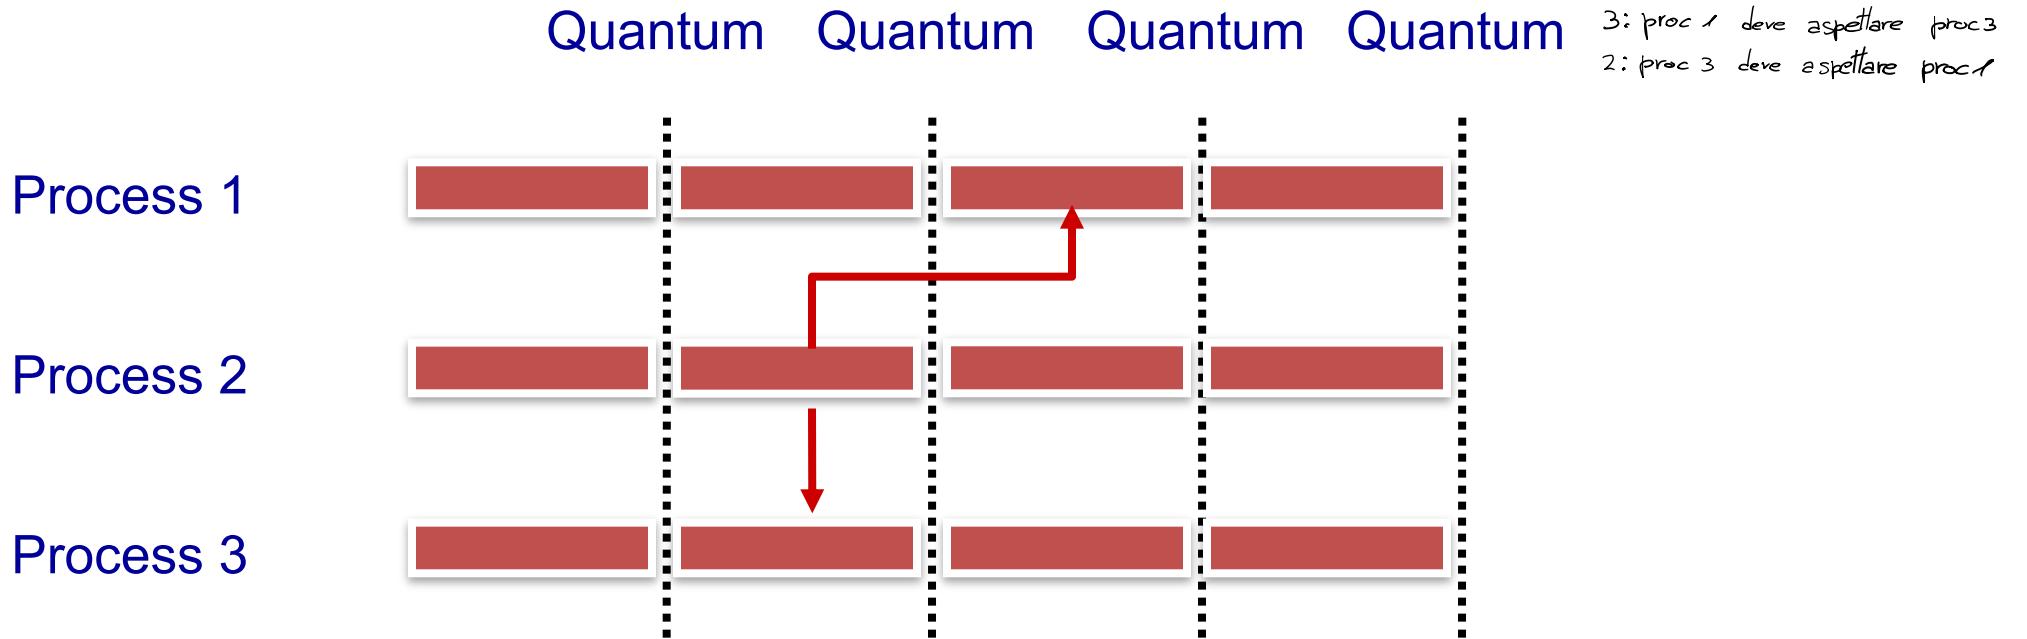
Coding Styles

- A coding style defines how time and data are related together, according to designers' perception
- It is NOT a specific abstraction level
- More of a guideline, designers enjoy several degrees of freedom
- TLM 2.0 standard provides two coding styles
 - Loosely-timed (LT)
 - Approximately-timed (AT)
- It is possible to define other coding styles

Loosely-timed (LT)

- Sufficient timing details to boot an operating system and run multi-core systems
- Allows processes to run ahead of simulation time (temporal decoupling) → faster simulation
- Each transaction has two timing points
 - Start and end of transaction
- Loosely-timed models have a limited number of temporal dependencies and context switches → faster simulation

Loosely-timed



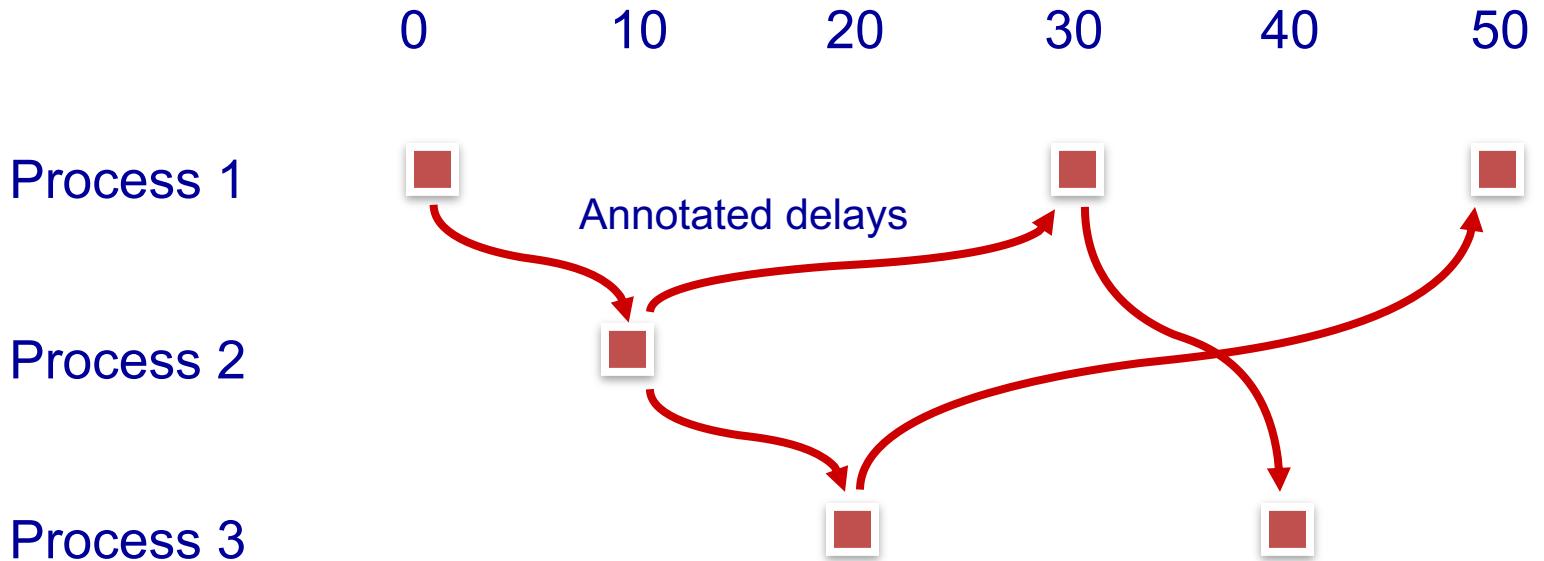
*Each process runs ahead up to quantum boundary
`sc_time_stamp()` advances in multiples of the quantum
 Deterministic communication requires explicit synchronization*

Approximately-timed (AT)

- Appropriate for architectural exploration and performance analysis
- Transactions are broken down into multiple phases
 - For example, four phases in the base protocol
 - Beginning of the request (BEGIN_REQ)
 - End of the request (END_REQ)
 - Beginning of the response (BEGIN_RESP)
 - End of the response (END_RESP)
 - It is possible to create specific protocols with a different number of phases
- Processes run in lock-step with simulation time
- Timing accuracy is needed, so process interactions are annotated with specific delays

Se continuo a sincronizzare initiator e target alla fine avro' il tempo esatto

Approximately-timed



*Each process is synchronized with SystemC scheduler
 Delays can be accurate or approximate*

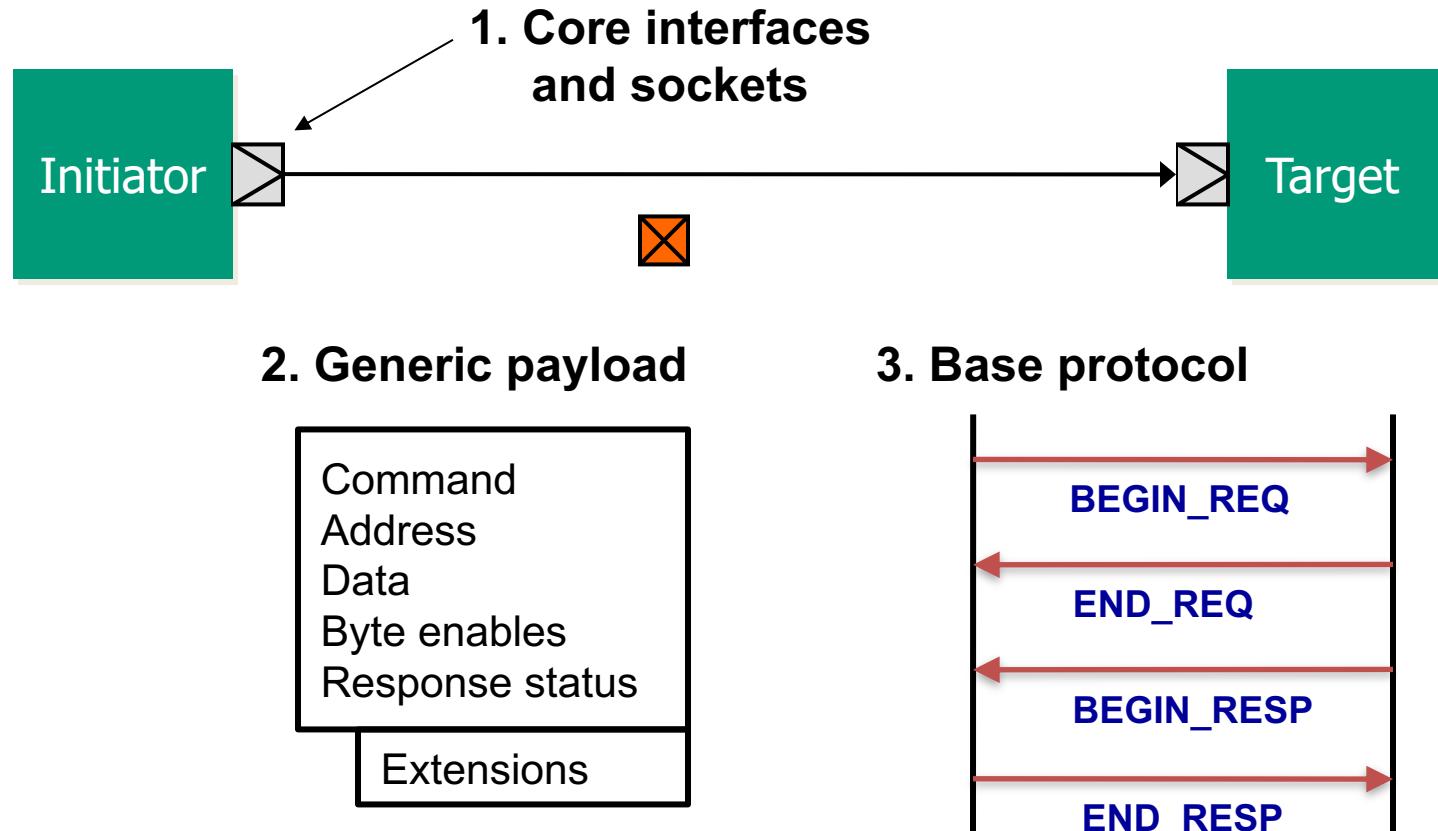
Untimed

- The notion of simulation time is unnecessary, so it is not taken into account
- TLM 2.0 does not explicitly support this coding style, because all contemporary bus-based systems require some notion of time to model software running on embedded processors
- An untimed model may be obtained by removing timing annotation, quantum keeping and temporal decoupling from a loosely-timed model

Interfaces

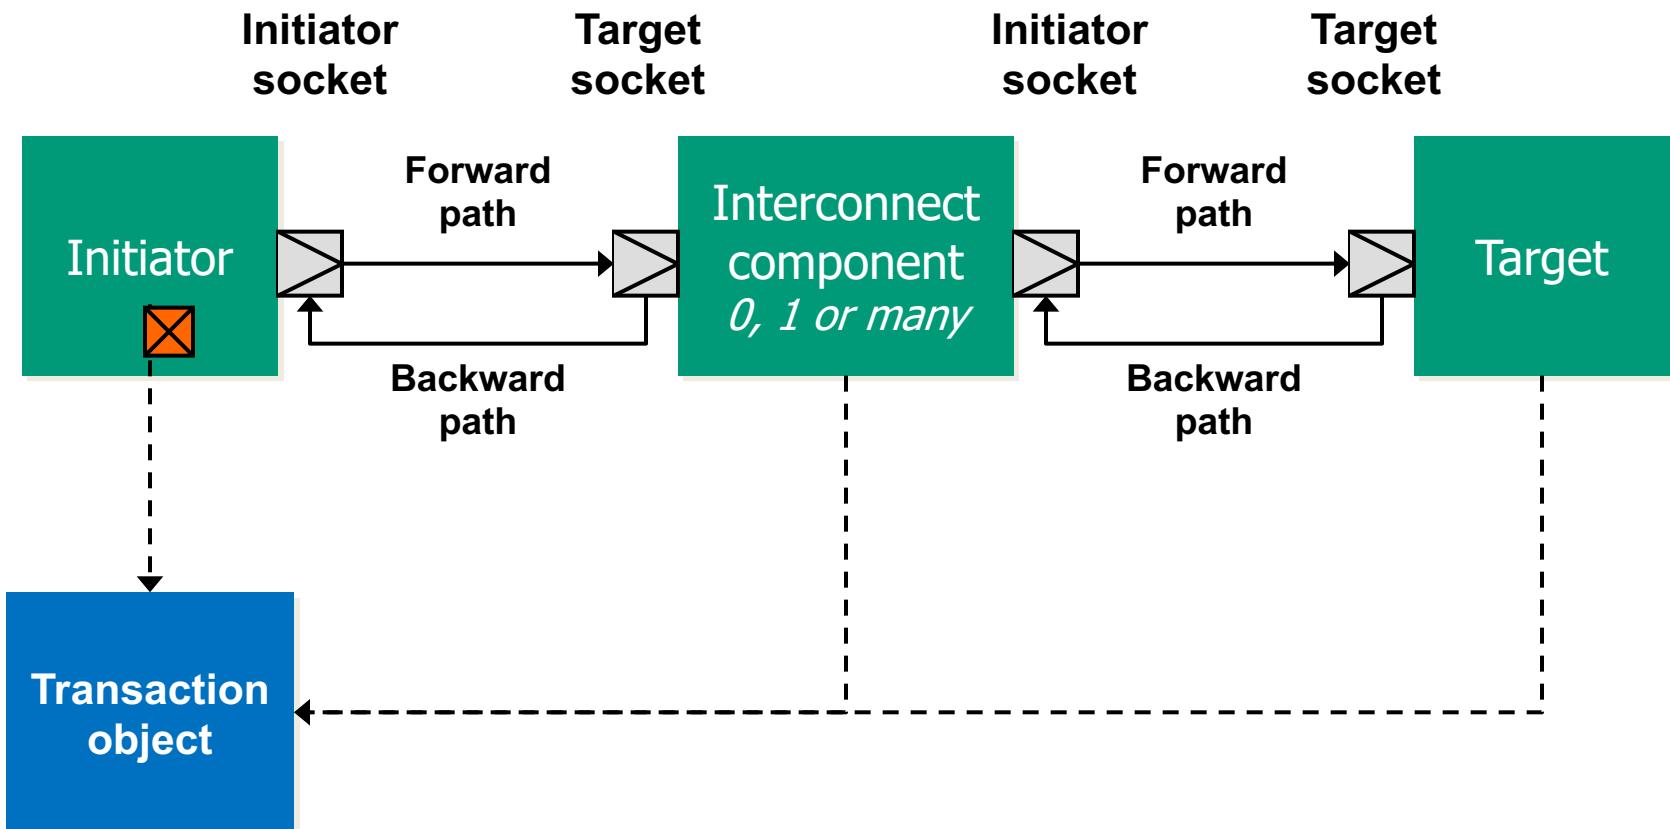
- Blocking transport interface
 - Intended to support the loosely-timed coding style
- Non-blocking transport interfaces
 - Intended to support the approximately-timed coding style
- Direct memory interface (DMI)
 - Allows the initiator to access an area of memory owned by the target by using a direct pointer
- Debug interface
 - Allows the initiator to communicate with the target without delays, waits, event notifications

Interoperability Layer



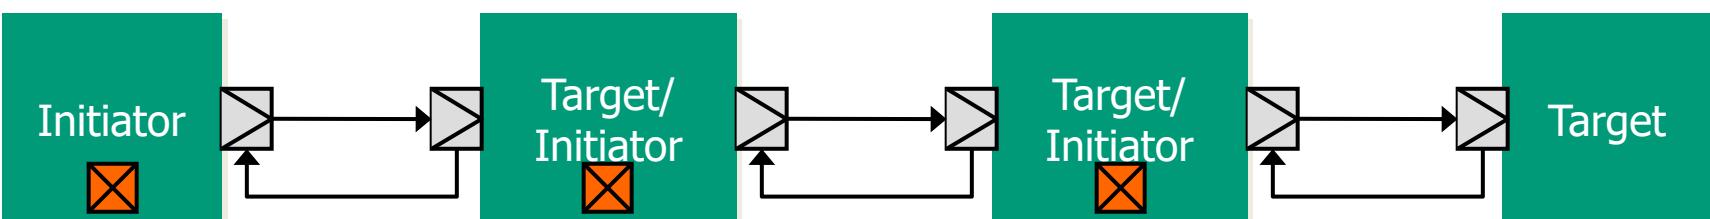
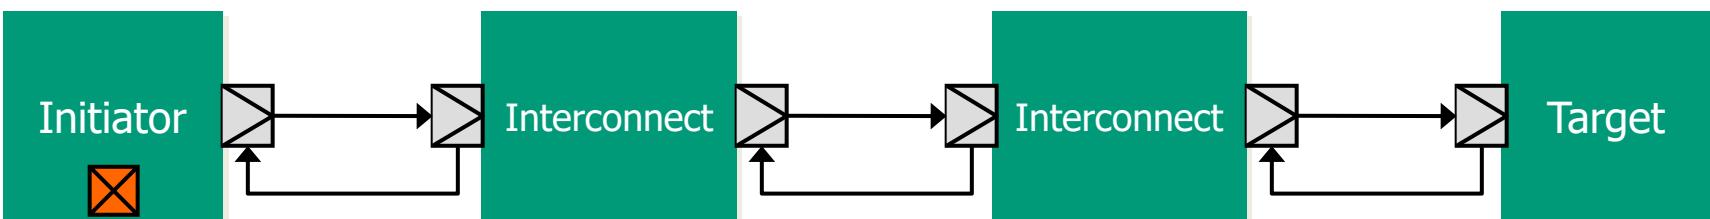
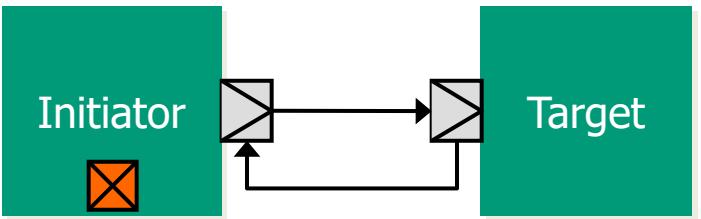
Maximal interoperability for memory-mapped bus models

Initiators and Targets



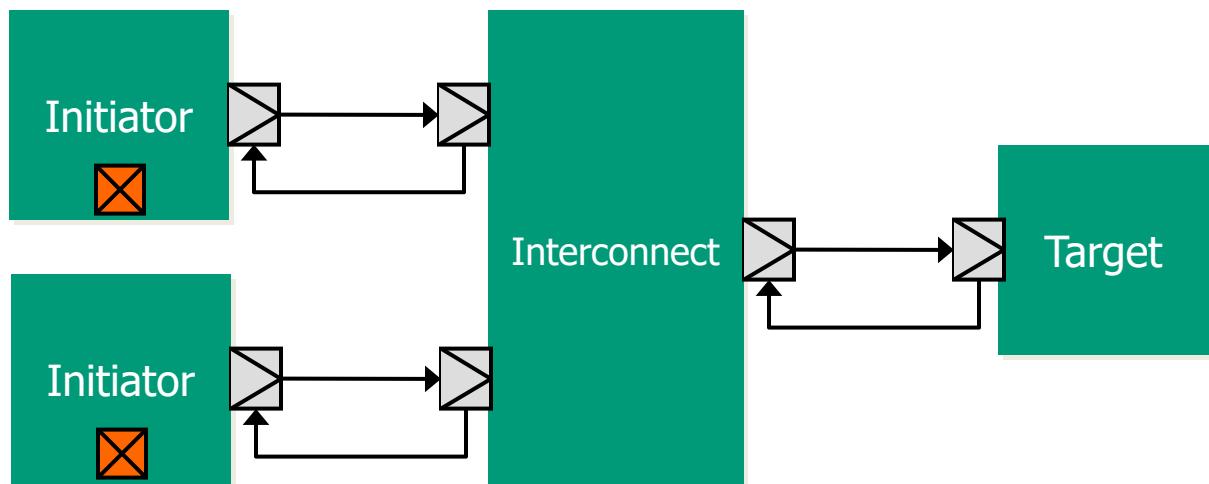
References to a single transaction object are passed along the forward and backward paths

TLM-2 Connectivity



Transaction memory management needed

Convergent Paths



Blocking versus Non-blocking Transport

- **Blocking transport interface**
 - Includes timing annotation
 - Typically used with loosely-timed coding style
 - Forward path only
- **Non-blocking transport interface**
 - Includes timing annotation and transaction phases
 - Typically used with approximately-timed coding style
 - Called on forward and backward paths
- Share the same transaction type for interoperability
- Unified interface and sockets – can be mixed

Blocking Interface

- Appropriate where an initiator wishes to complete a transaction with a target in a single function call
 - Two timing points
 - Call to and return from the blocking transport function
 - It only uses the forward path from initiator to target

Transaction type

↓

```
template < typename TRANS = tlm_generic_payload >

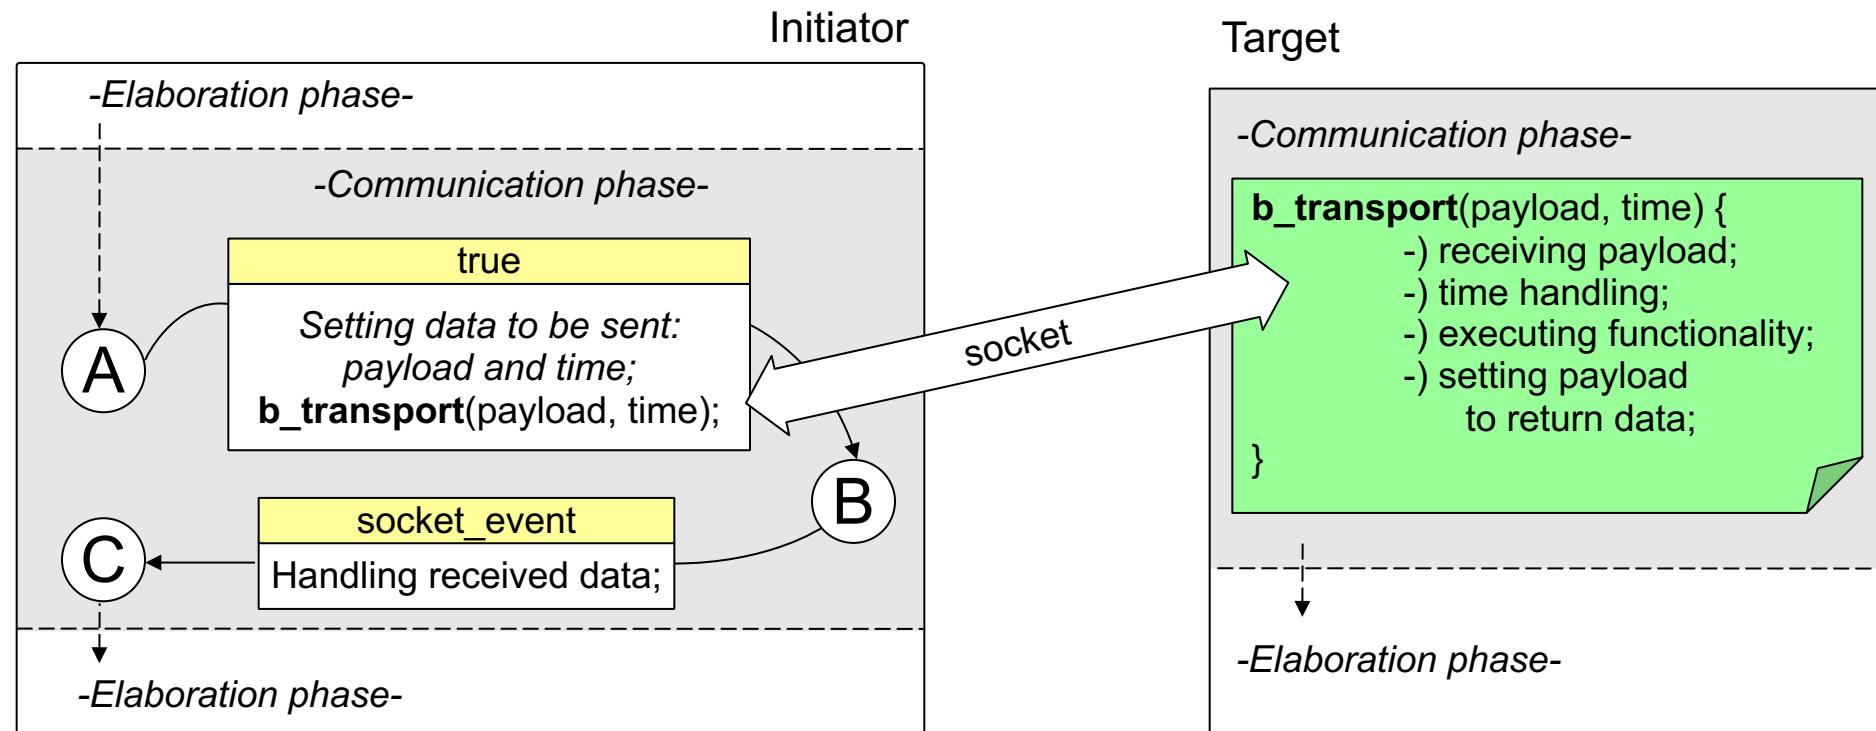
class tlm_blocking_transport_if : public virtual sc_core::sc_interface {
public:
    virtual void b_transport ( TRANS& trans , sc_core::sc_time& t ) = 0;
};
```

Transaction object

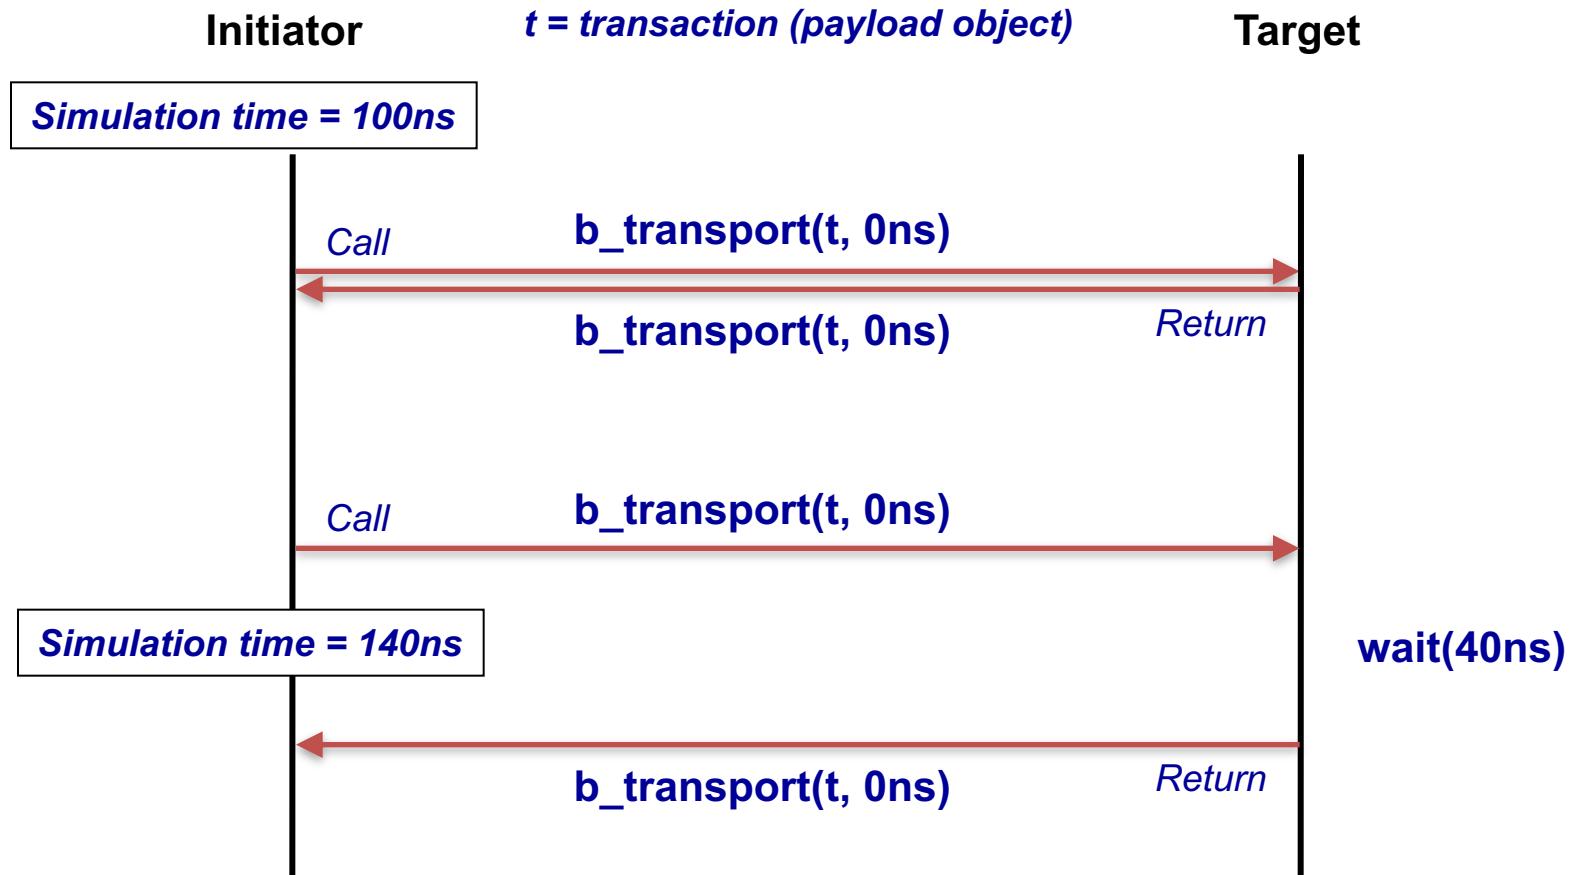
↑

Timing annotation

EFSM Model of a LT-based Protocol



Blocking Interface Example



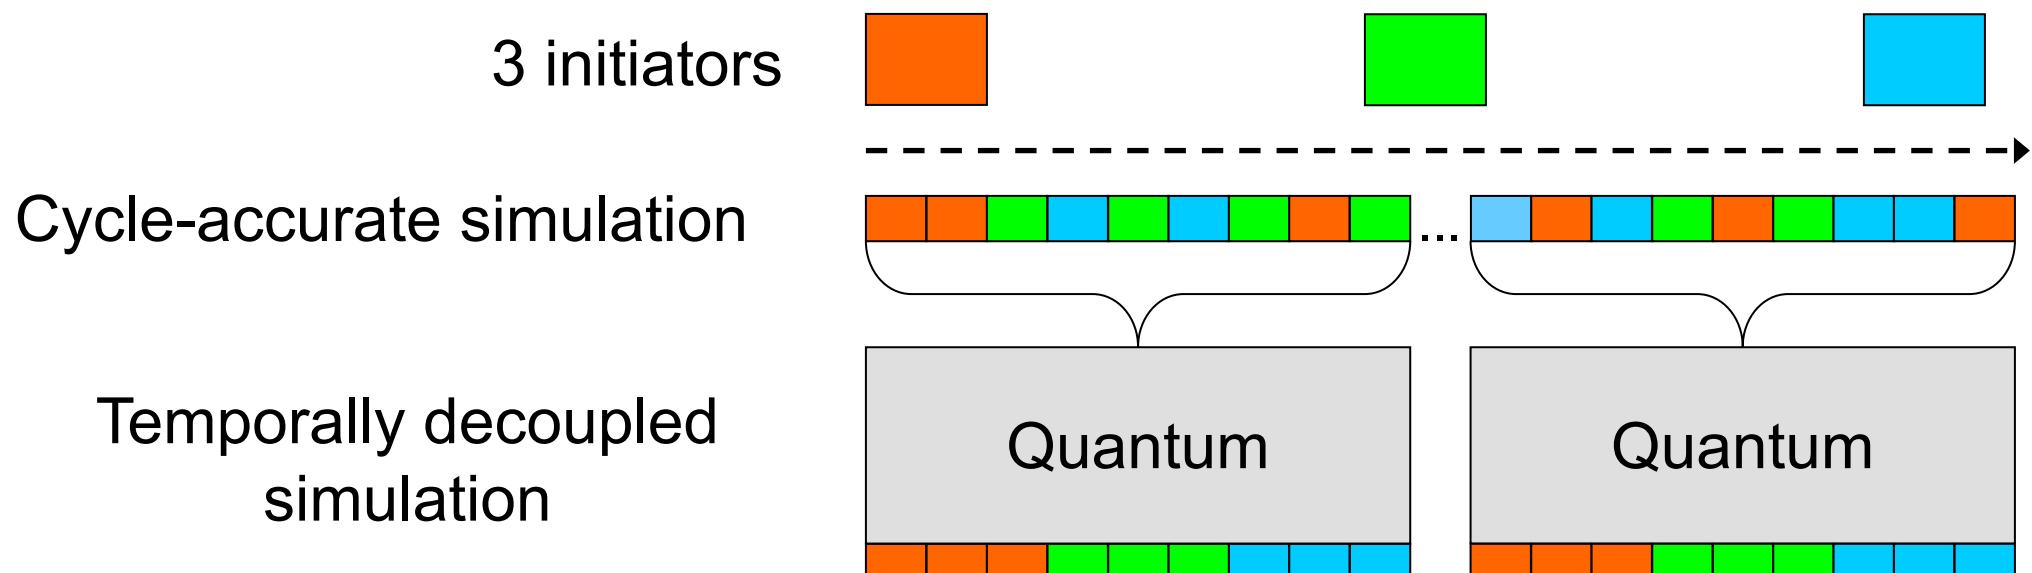
Initiator is blocked until return from b_transport

Temporal Decoupling

- A temporally decoupled initiator may run at a notional local time in advance of the current simulation time
 - in this case it should pass a non-zero value for the time argument to the `b_transport()` primitive
- Each process is allowed to run for a certain time slice (time quantum) before switching to the next process
- A temporally decoupled initiator will continue to advance local time until the time quantum is exceeded or it reaches an explicit synchronization point
- At that point, the initiator is obliged to synchronize by suspending execution
- Simulation time cannot advance until the initiator thread yields
- The quantum keeper is responsible for enforcing quantum timing

Time Quantum

- Each process runs ahead up to quantum boundary
- Simulation time advances in multiples of the quantum
- Fewer context switches → faster simulation, but also less time accuracy
- Deterministic communication requires explicit synchronization



Quantum Keeper Example (I)

```

struct Initiator: sc_module {
    tlm_utils::tlm_quantumkeeper m_qk; ← The quantum keeper object

    SC_CTOR(Initiator) : init_socket("init_socket") {
        m_qk.set_global_quantum( sc_time(500, SC_NS) ); ← Set the global quantum
        m_qk.reset(); ← Reset the local quantum
    }
    void thread() { ...
        for (int i = 0; i < RUN_LENGTH; i += 4) {
            local_time = m_qk.get_local_time(); ← Retrieve current local time
            init_socket->b_transport( trans, local_time );
            cout << "After b_transport - time = " << sc_time_stamp()
                << " + " << local_time << endl;
            m_qk.set( local_time ); ← Time consumed by transport
            m_qk.inc( sc_time(100, SC_NS) ); ← Further time consumed by initiator
            cout << "Initiator completed - time = " << sc_time_stamp()
                << " + " << m_qk.get_local_time() << endl;
            if ( m_qk.need_sync() ) { ← Check local time against quantum
                cout << "### SYNCHRONIZING ###" << endl;
                m_qk.sync(); ← Synchronize if necessary
            }
        }
    };
}
  
```

i = 0 i = 1

0 + 150 ns	0 + 400 ns
0 + 250 ns	0 + 500 ns
	500 + 0 ns

Simulation time } Local time }

Non-blocking interfaces (I)

- Appropriate where it is desired to model the detailed sequence of interactions between initiator and target during the course of each transaction
 - Multiple function calls to complete a transaction
- Multiple phases and timing points
 - Each phase transition is associated with a timing point
- It uses both the forward and the backward path
- Two distinct interfaces for use on opposite paths
 - `tlm_fw_nonblocking_transport_if`
 - `tlm_bw_nonblocking_transport_if`

Non-blocking interfaces (II)

```

enum tlm_sync_enum { TLM_ACCEPTED, TLM_UPDATED, TLM_COMPLETED };

template < typename TRANS = tlm_generic_payload,           ← Transaction type
          typename PHASE = tlm_phase>                      ← Phase type

class tlm_fw_nonblocking_transport_if : public virtual sc_core::sc_interface {
public:
    virtual tlm_sync_enum nb_transport_fw(      TRANS& trans,
                                              PHASE& phase,
                                              sc_core::sc_time& t ) = 0;
};

class tlm_bw_nonblocking_transport_if : public virtual sc_core::sc_interface {
public:
    virtual tlm_sync_enum nb_transport_bw(      TRANS& trans,   ← Transaction object
                                              PHASE& phase,   ← Phase object
                                              sc_core::sc_time& t ) = 0;
};

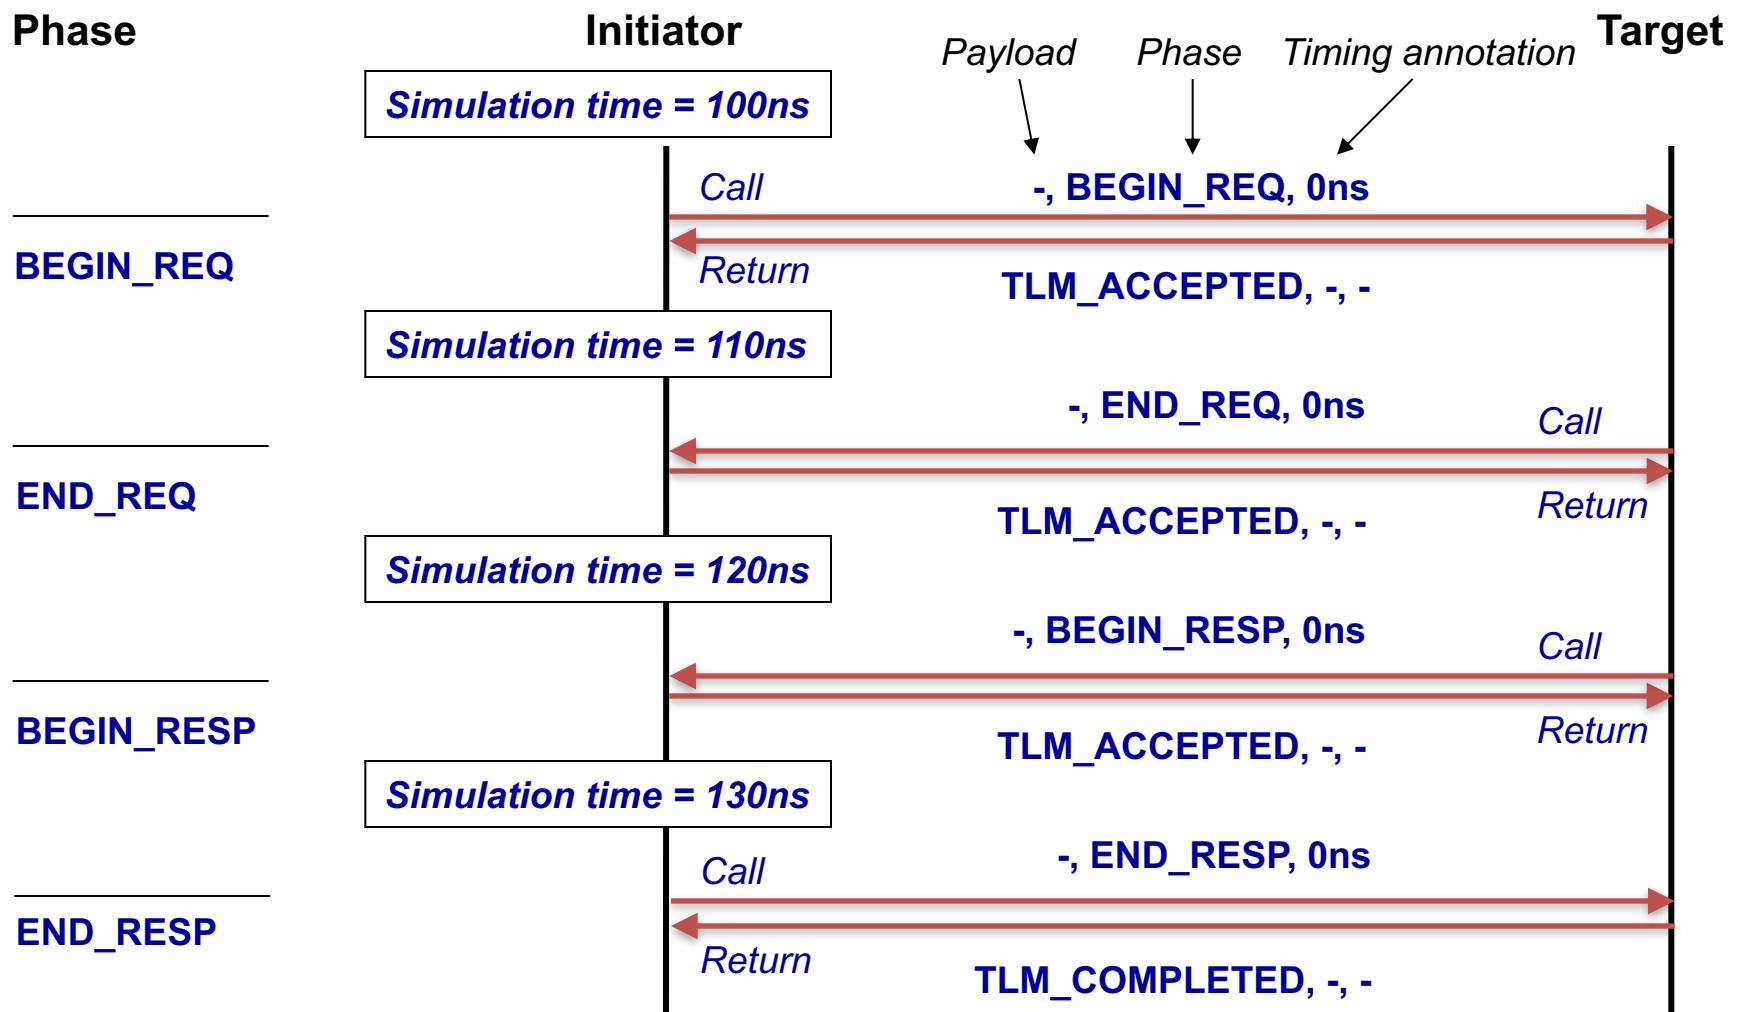
```

\uparrow
Timing annotation

tlm_sync_enum overview

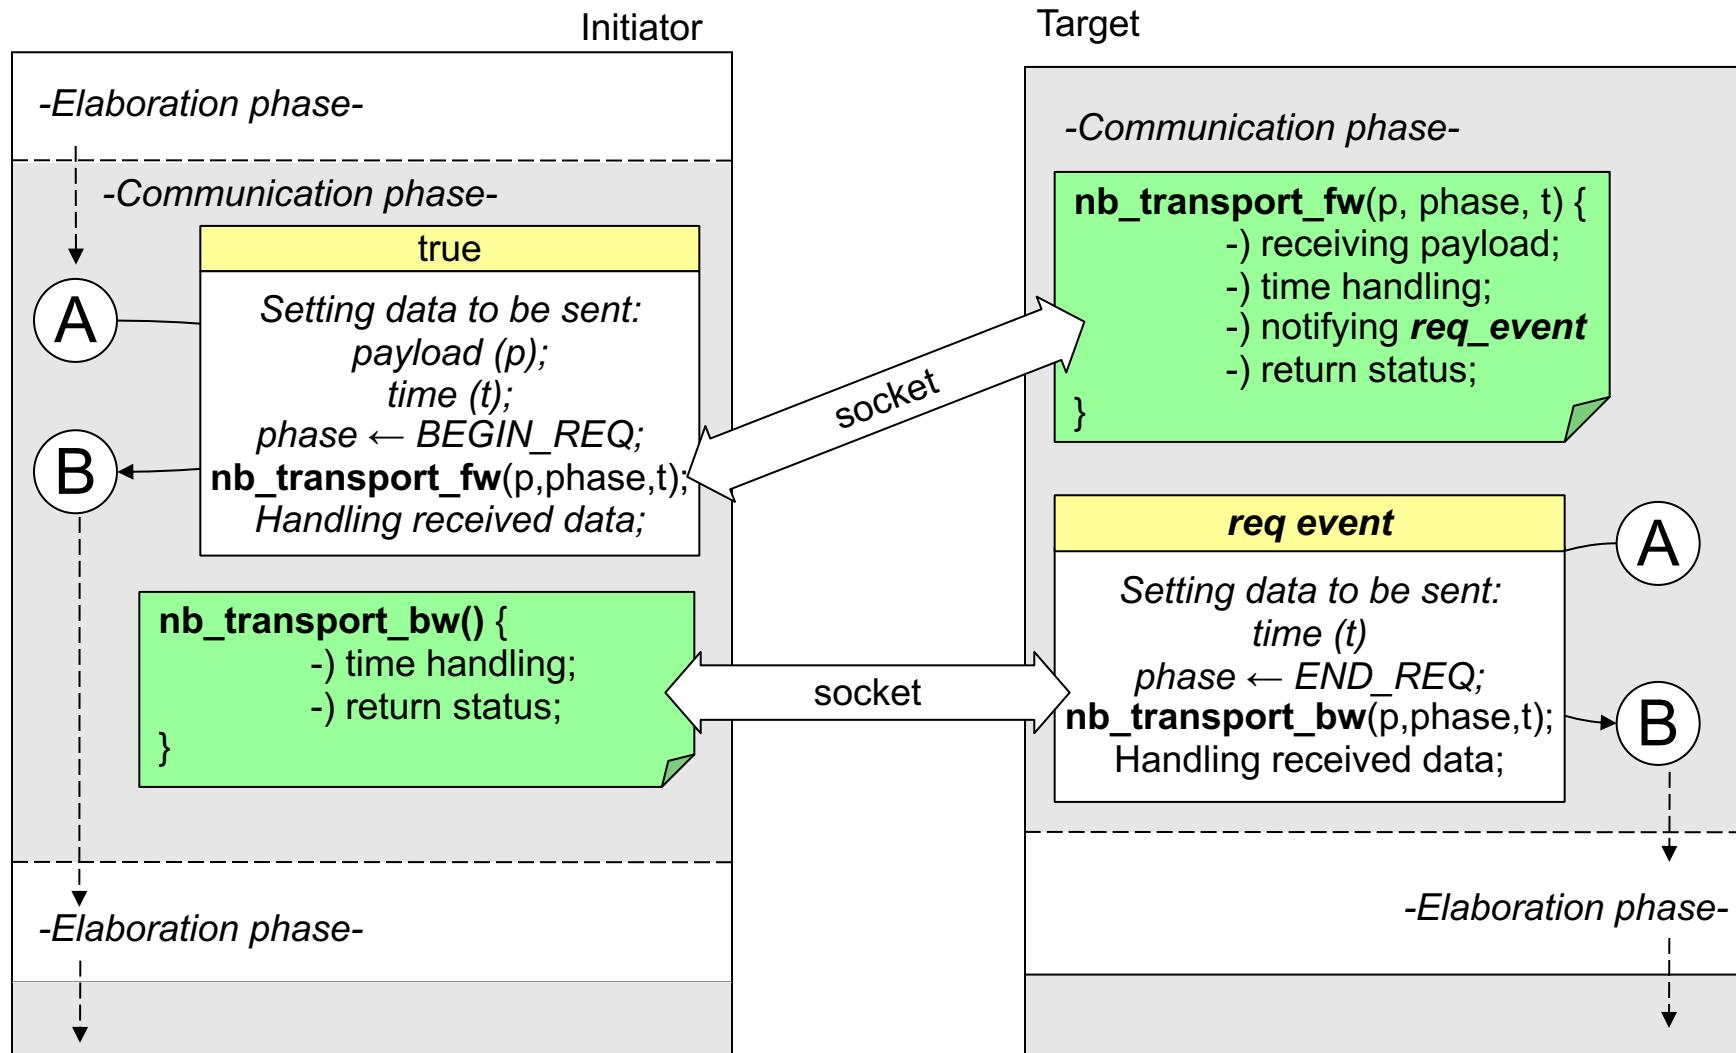
- **TLM_ACCEPTED**
 - The callee shall not have modified the transaction object, the phase or the time argument during the call
- **TLM_UPDATED**
 - The callee has updated the transaction object
 - It may have modified the phase argument
 - It may have increased the value of the time argument
- **TLM_COMPLETED**
 - The callee has updated the transaction object
 - It may have increased the value of the time argument
 - The transaction is complete

Non-blocking Interface Example (I)

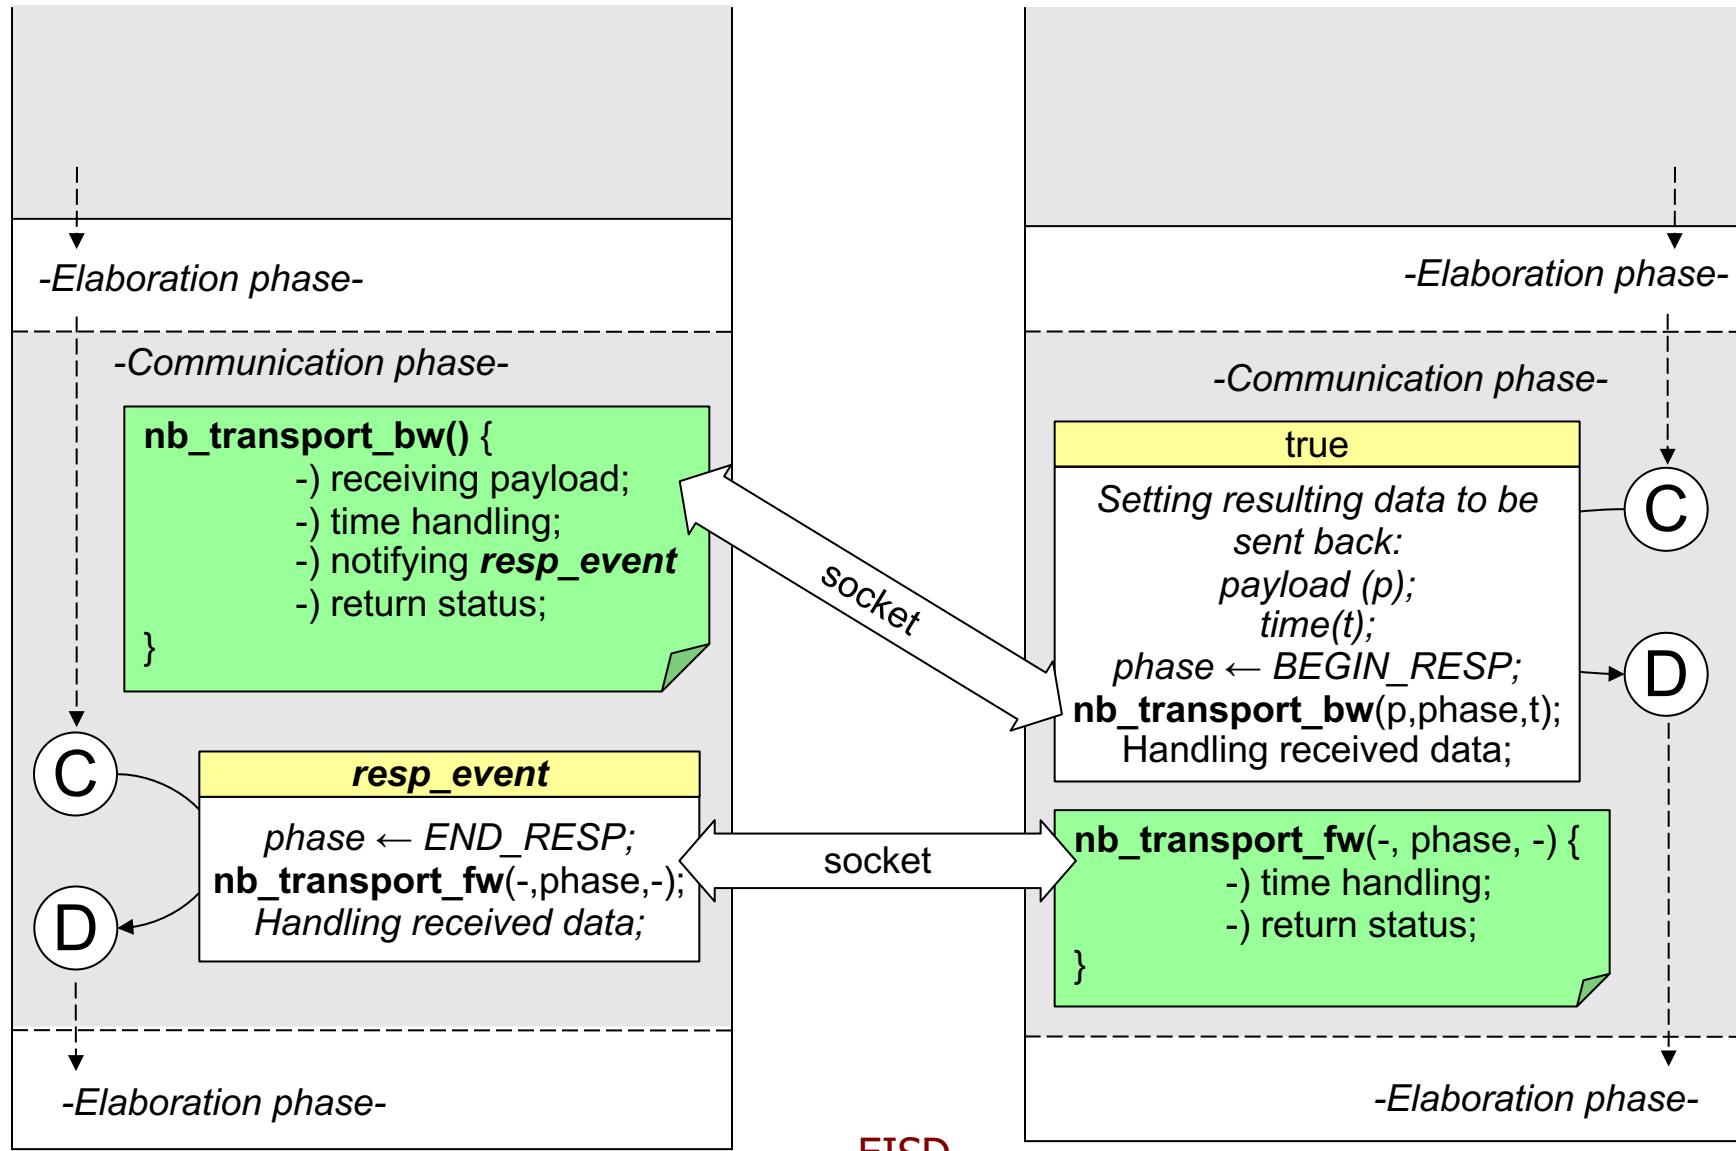


Transaction accepted now, caller asked to wait

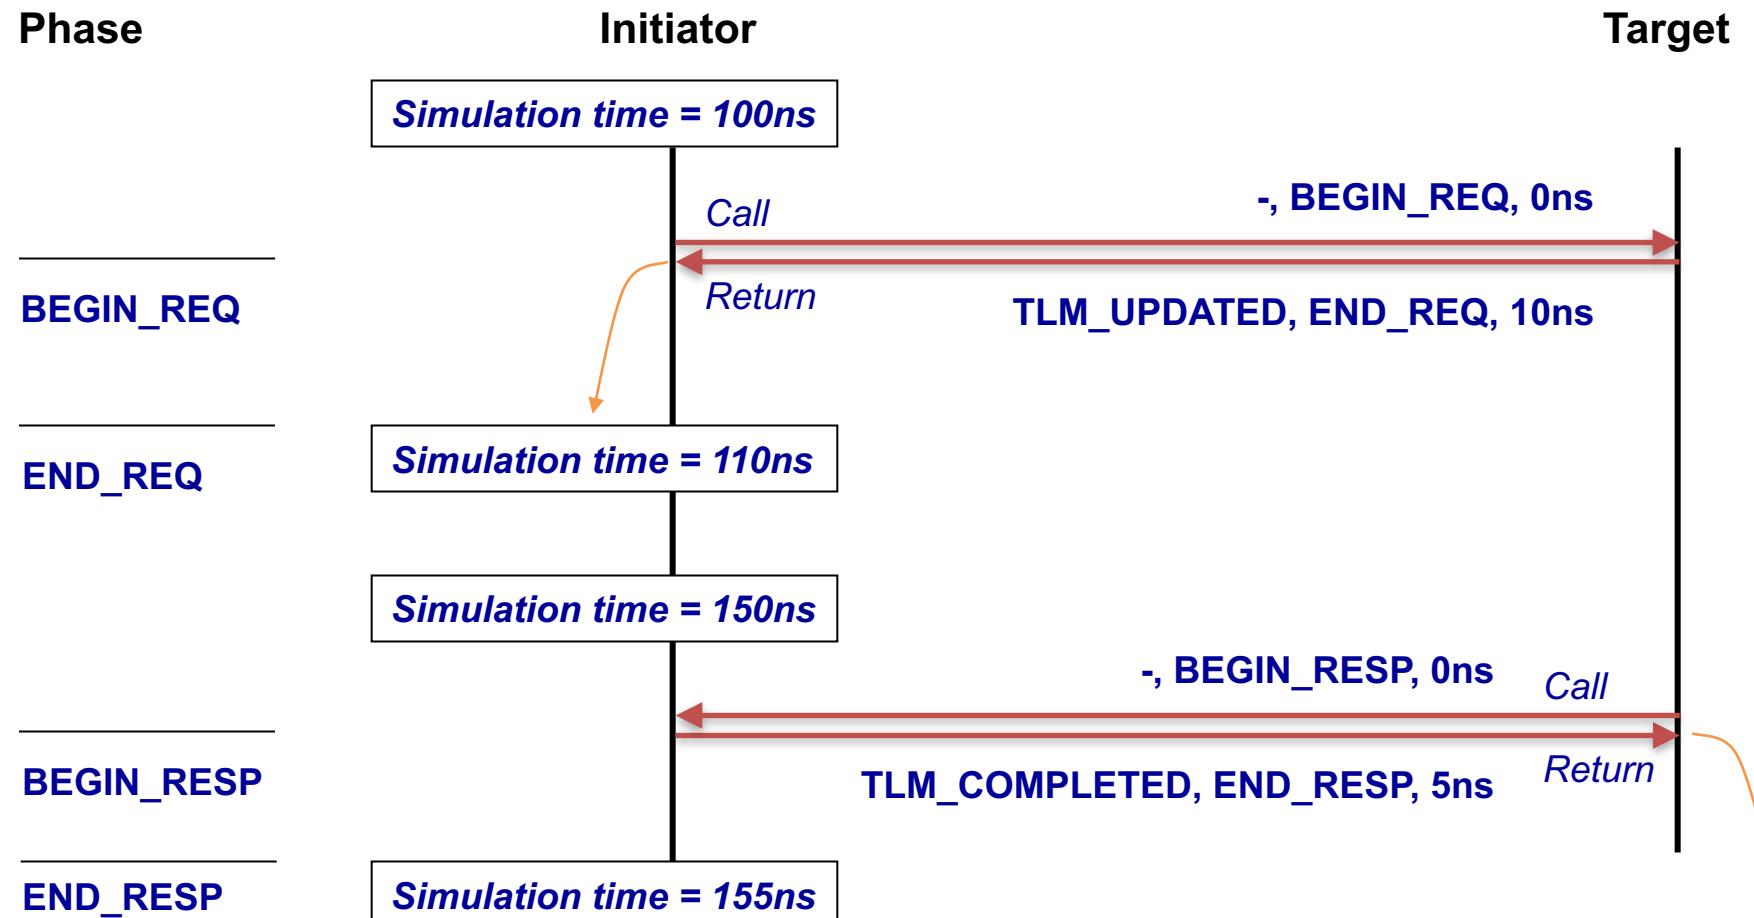
EFSM Model of an AT-based Protocol (I)



EFSM Model of an AT-based Protocol (II)



Non-blocking interface example (II)



Callee annotates delay to next phase transition, caller waits

The Generic Payload

Attribute	Type	Modifiable?	
Command	<code>tlm_command</code>	No	
Address	<code>uint64</code>	Interconnect only	
Data pointer	<code>unsigned char*</code>	No (array – yes)	<i>Array owned by initiator</i>
Data length	<code>unsigned int</code>	No	
Byte enable pointer	<code>unsigned char*</code>	No (array – yes)	<i>Array owned by initiator</i>
Byte enable length	<code>unsigned int</code>	No	
Streaming width	<code>unsigned int</code>	No	
DMI hint	<code>bool</code>	Yes	
Response status	<code>tlm_response_status</code>	Target only	
Extensions	<code>(tlm_extension_base*)[]</code>	Yes	<i>Consider memory management</i>

↳ Puntatori ad altre strutture dati.

Response Status

enum tlm_response_status	Meaning
TLM_OK_RESPONSE	Successful
TLM_INCOMPLETE_RESPONSE	Transaction not delivered to target. (Default)
TLM_ADDRESS_ERROR_RESPONSE	Unable to act on address
TLM_COMMAND_ERROR_RESPONSE	Unable to execute command
TLM_BURST_ERROR_RESPONSE	Unable to act on data length or streaming width
TLM_BYTE_ENABLE_ERROR_RESPONSE	Unable to act on byte enable
TLM_GENERIC_ERROR_RESPONSE	Any other error

Key Payload Fields

- Command
 - indicates the type of operation (read, write or none) requested to the target
- Address
 - indicates the address in target's memory of data to be read or written
- Data pointer
 - pointer to the data array
 - in case of a read operation, the target will copy the content of a local array in the data array
 - in case of a write operation, the target will copy the content of the data array in a local array
- Data length
 - provides the number of bytes to be stored in or read from the data array
- Response status
 - indicates the outcome of the operation performed by the target

Generic payload extensions

- The generic payload provides an off-the-shelf general purpose payload for abstract bus modeling
 - It is aimed at modeling memory-mapped buses
- To model other specific protocols, the generic payload includes an extension mechanism, so that specialized attributes can be added
 - Generic payload has an array-of-pointers to extensions
 - One pointer per extension type
 - Every transaction can potentially carry every extension type
 - This mechanism allows greater flexibility

System Verification

Step iniziale di verifica tra descrizione e prima versione del sistema

- Validation:
 - Does the design correctly work?
 - Are there implementation design errors?
 - No reference models!
- Verification:
 - Refinement (or abstraction) step verification:
ascertaining that system functionality is kept unaltered throughout its representation hierarchy
 - Golden Model as reference

⇒ SISTEMA IDEALE PER COMPARARE IL MIO SISTEMA ED ESSERE SICURO CHE FUNZIONI

System Verification Techniques

Dynamic verification
(simulation)

- ↑ High performance
- ↑ Advanced commercial tools
- ↓ No-exhaustive technique

Static verification
(formal)

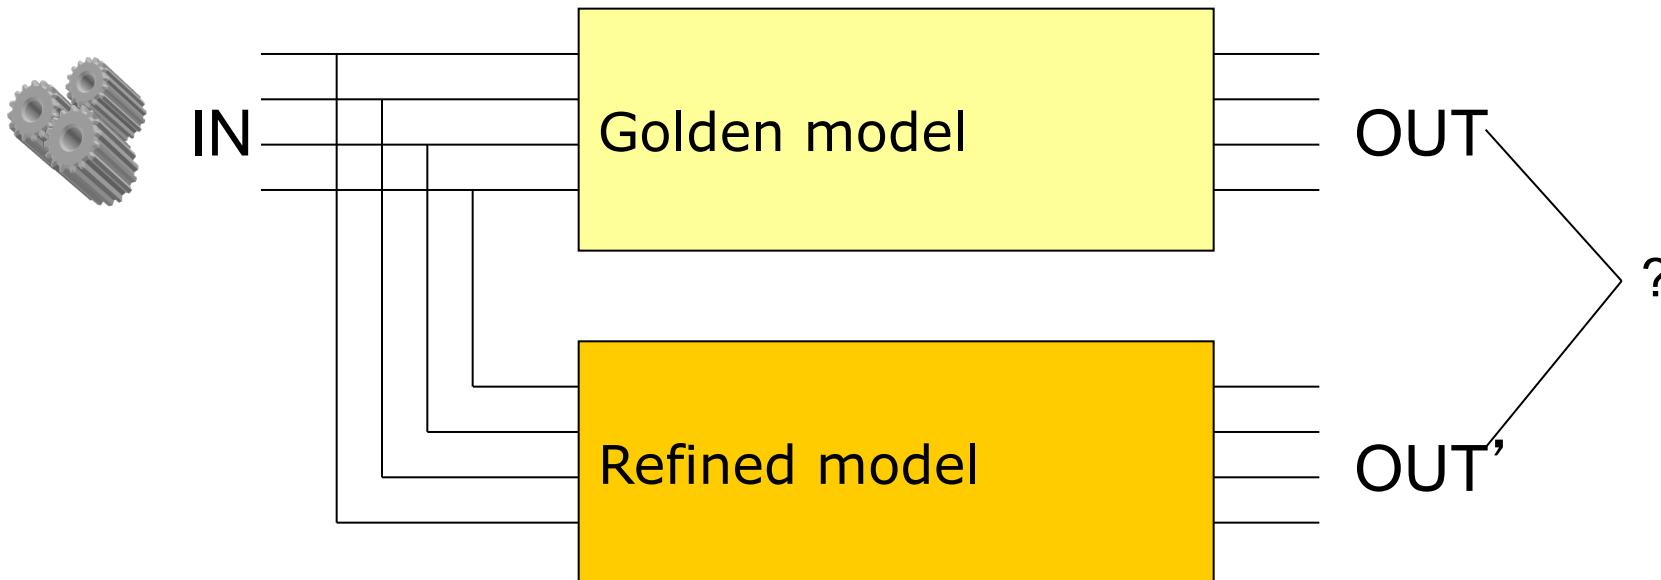
- ↑ Quality
- ↑ Exhaustive technique
- ↓ Long verification time
- ↓ For small size designs

Semi-formal verification



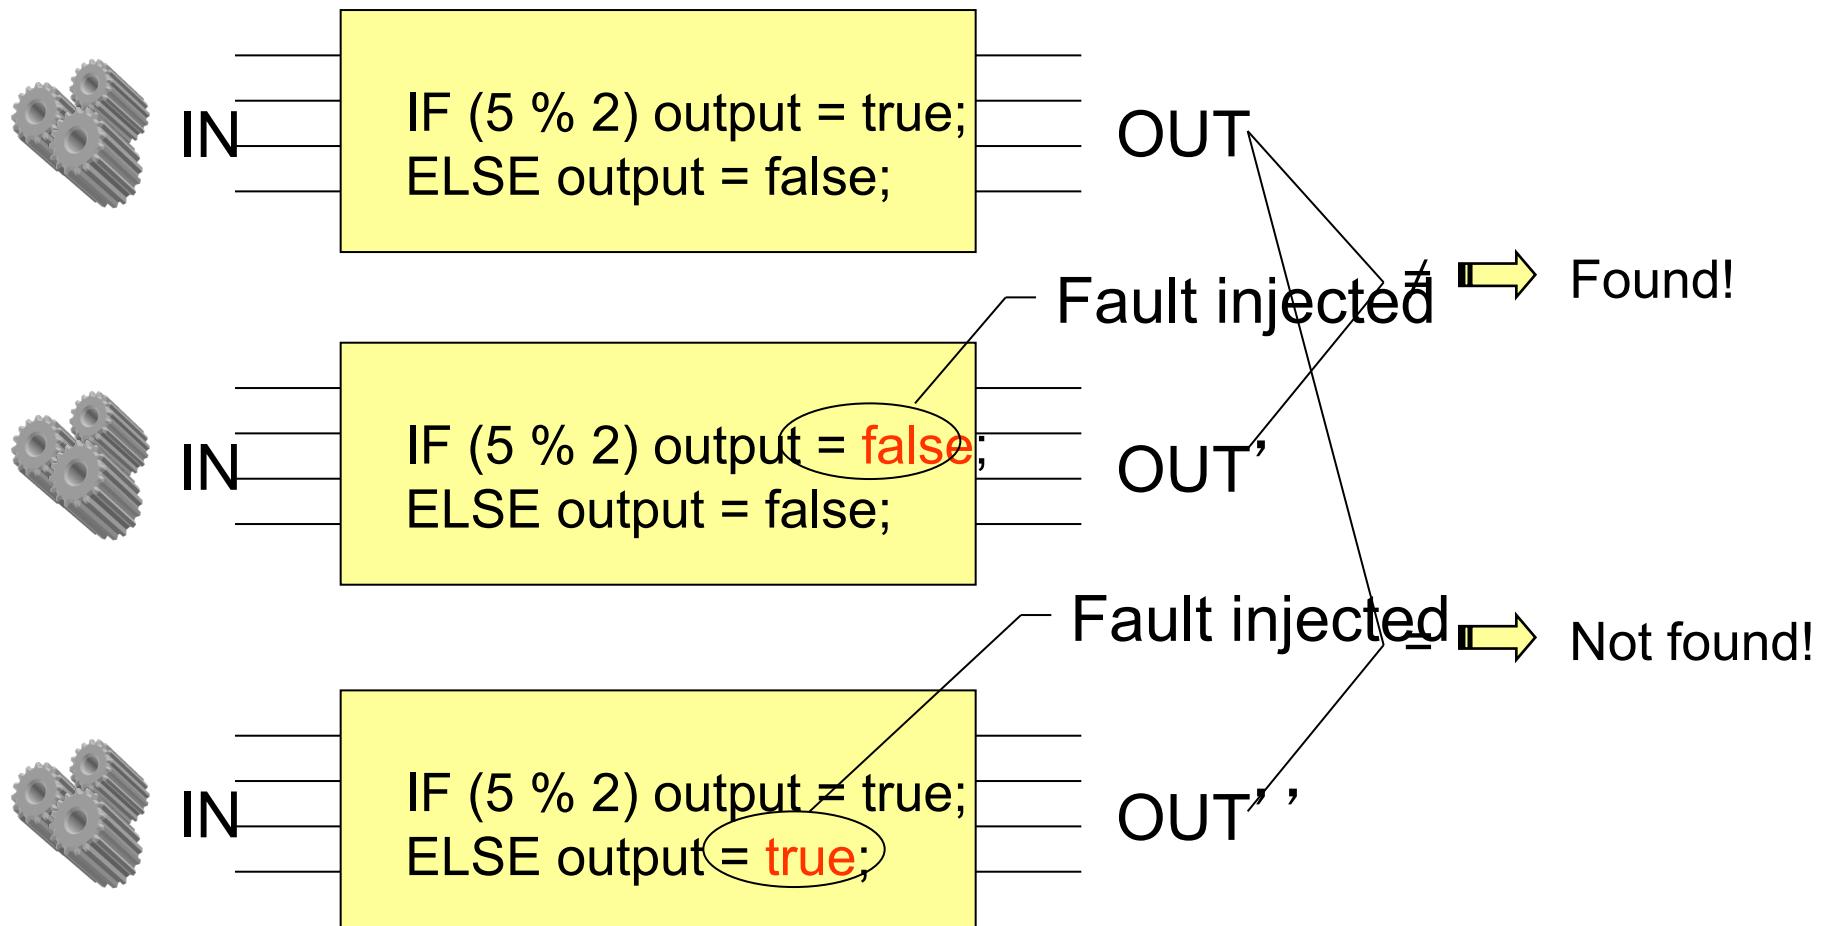
STATO DELL'ARTE ATTUALE

Dynamic Verification (simulation)



- ↑ System level and block level verification
- ↑ Easiest
- ↑ Suitable for TLM & RTL
- ↓ Results quality

Dynamic Verification: Faults Injection



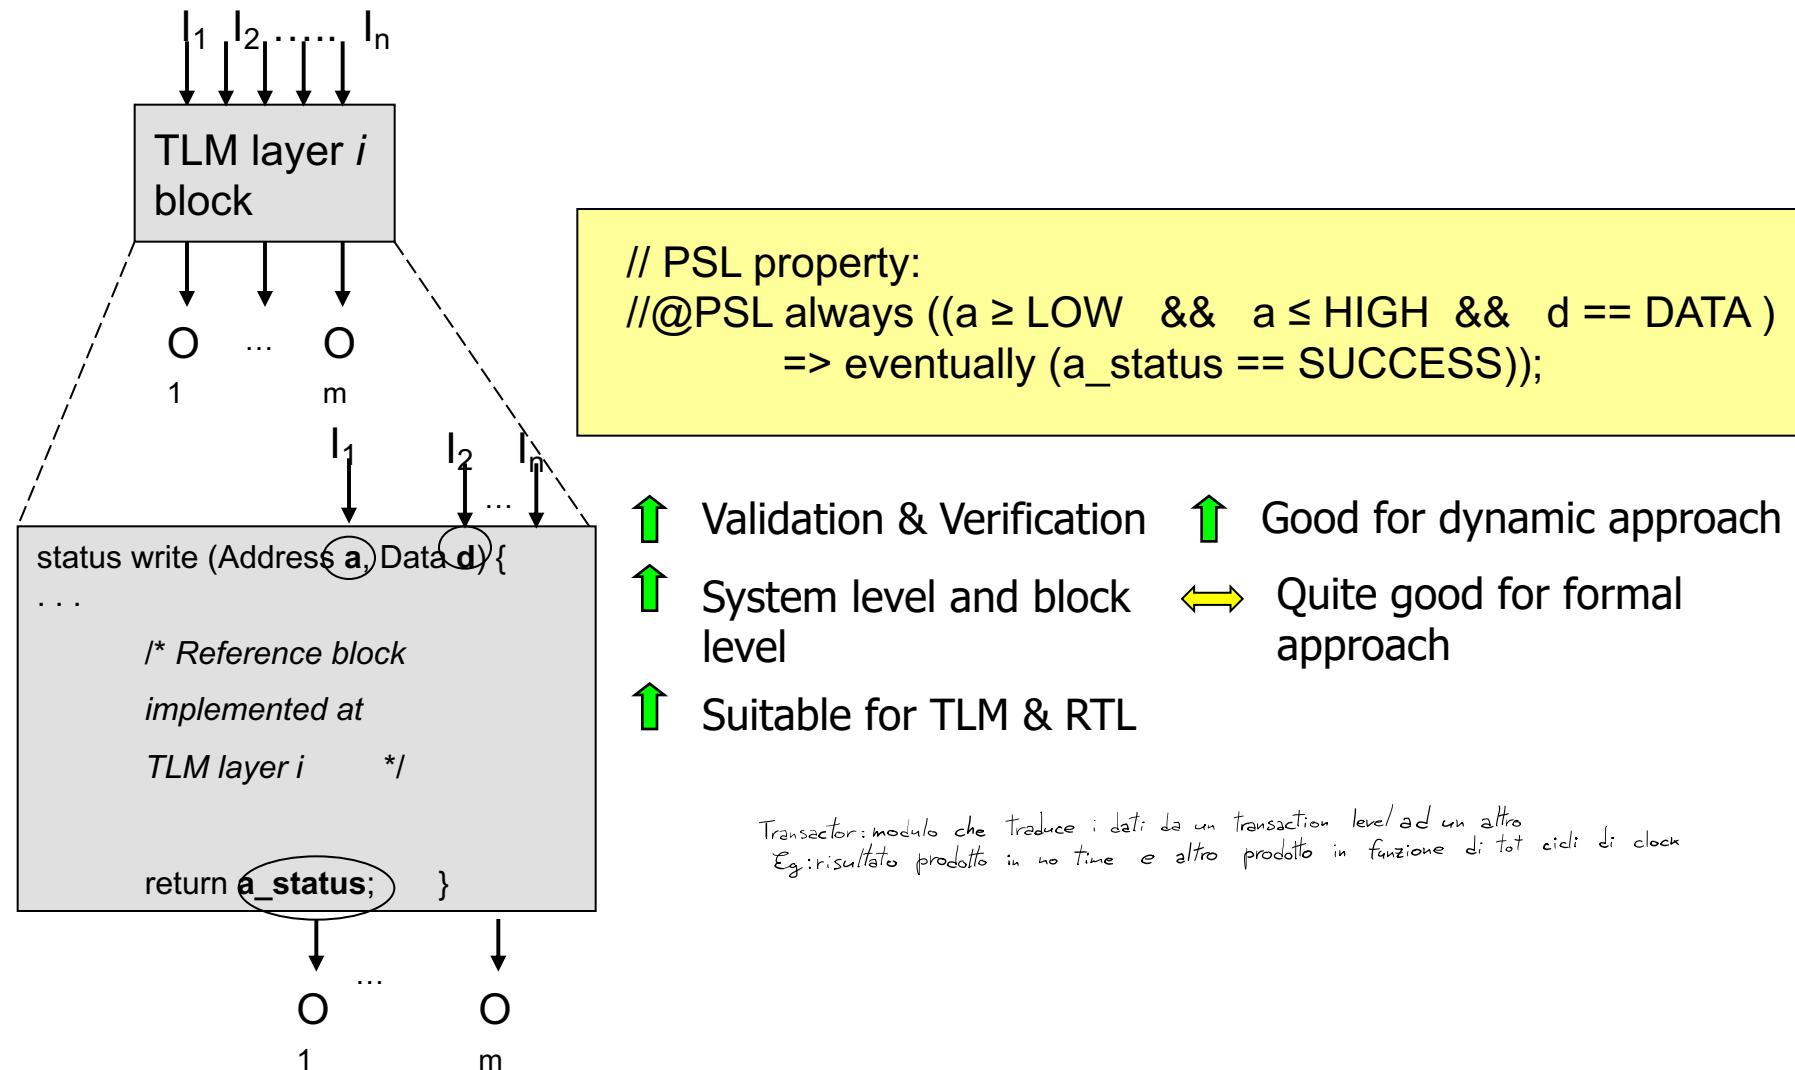
↔ Block level validation

↑ Results quality

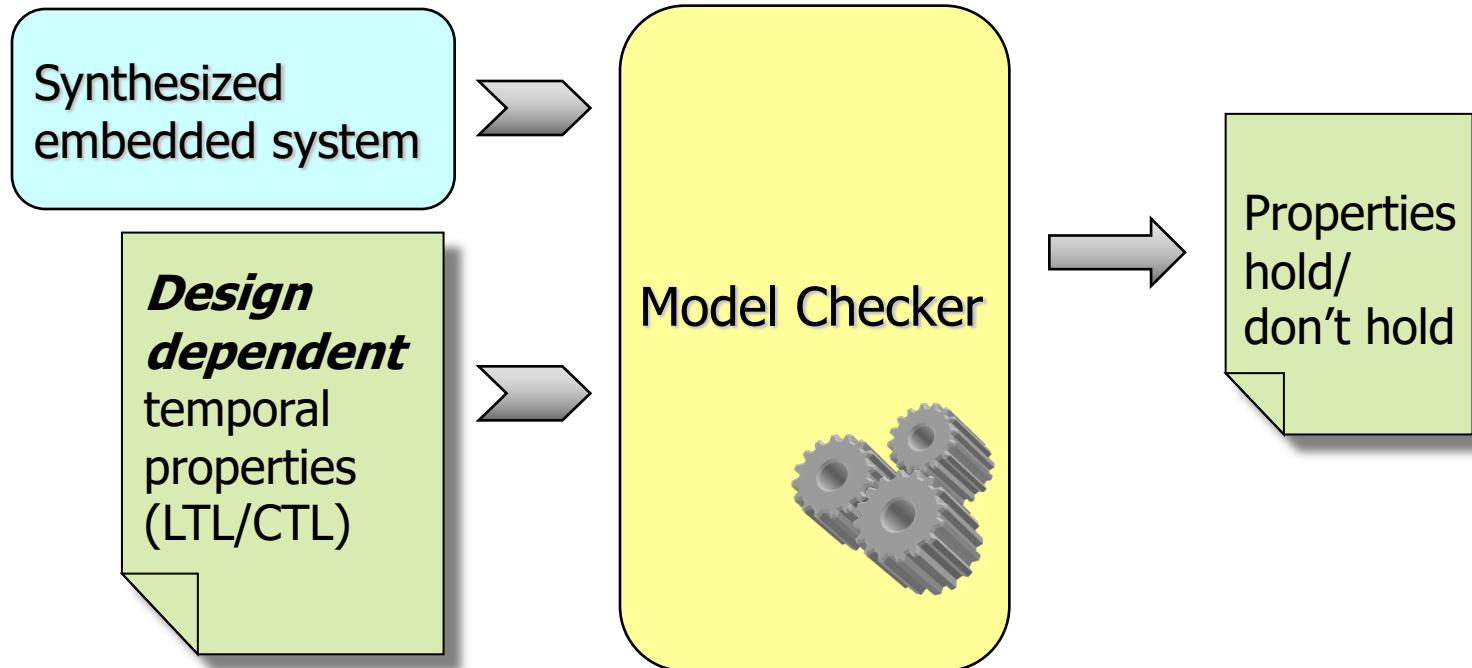
↔ Suitable for TLM level 1 & RTL

⬇ Synthesizable design required for faults injection

Property Specification Language (PSL)



Model Checking



↔ Block level validation

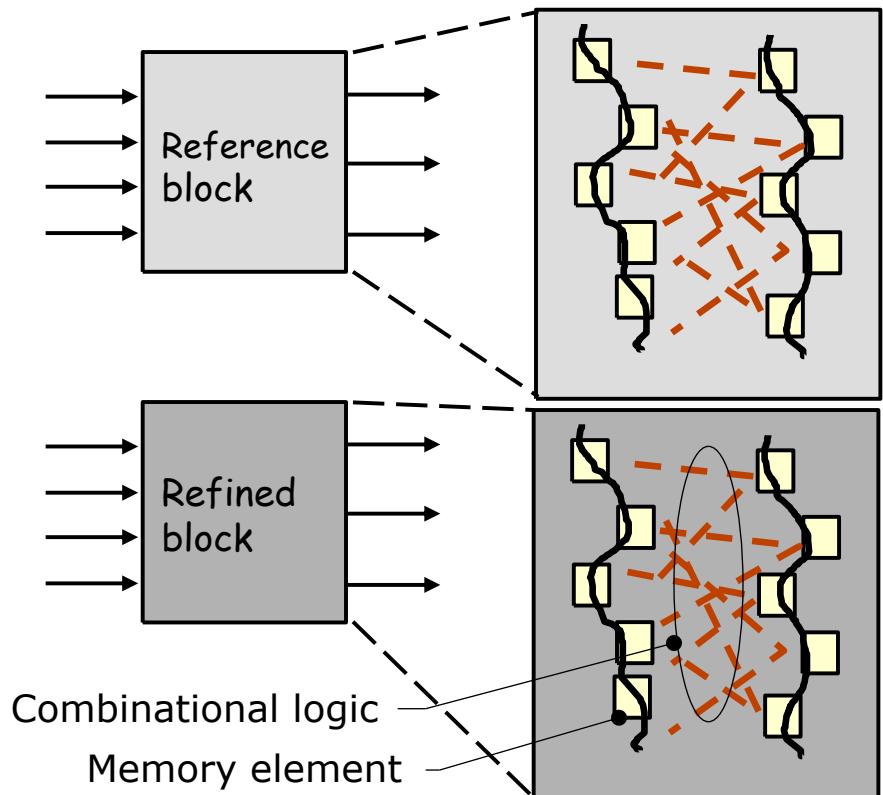
↑ Results quality

↔ Suitable for TLM level 1 & RTL

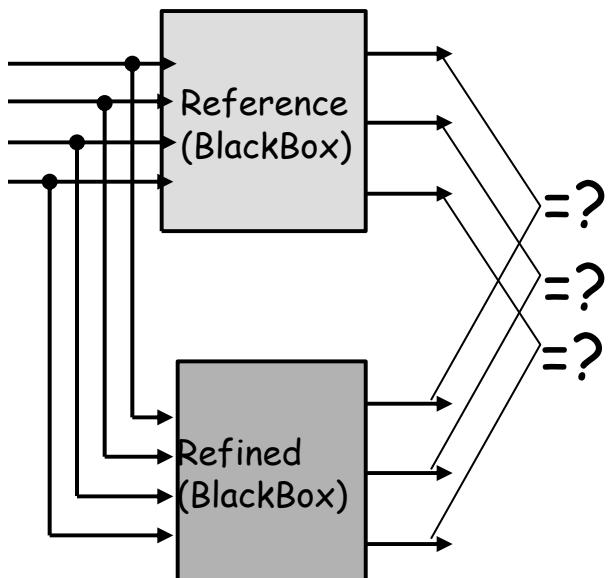
↓ Synthesizable design required

Equivalence Checking (EC)

Combinational EC



Sequential EC



Equivalence Checking (EC)

- Combinational Equivalence Checking

- | | |
|--|--|
|  Block level verification |  Synthesizable design required |
|  Results quality |  Suitable only for RTL vs Gate level verification |
| |  For small size designs |

- Sequential Equivalence Checking

- | | |
|---|--|
|  Block level verification |  Synthesizable design required |
|  Suitable for TLM & RTL |  For very small size design |
|  Results quality | |

Conclusions

- OSCI TLM 2.0 Standard (Accellera SystemC 2.3.1) provides:
 - Transport interfaces with timing annotations and phases
 - Loosely-timed coding style + temporal decoupling for simulation speed
 - Approximately-timed coding style for timing accuracy
 - Generic payload for memory-mapped bus modeling
 - Base protocol to enable interoperability between models
 - Interoperability allows reuse → reduction of production time and costs
 - Extensions for flexibility of modeling